

HDFS Task-1

Answer

1. What is the Namenode's URI and which file is it configured in?

URI : Uniform resource identifier (URI) in HDFS to locate a file.

it can be configured in core-site.xml file in "fs.default.name" name property

2. Where on a local file system will Namenode store its image and which file is it configured in?

/opt/hadoop-2.6.0/Hadoop-data/dfs/name

it is configured in "hdfs-site.xml" file

```
<property>
  <name>dfs.namenode.name.dir</name>
  <value>file:///opt/hadoop-2.6.0/hadoop_data/dfs/name</value>
</property>
```

3. Where on a local file system will Datanode store its blocks and which file is it configured in?

/opt/hadoop-2.6.0/Hadoop-data/dfs/name

it is configured in "hdfs-site.xml" file

```
<property>
  <name>dfs.datanode.data.dir</name>
  <value>file:///opt/hadoop-2.6.0/hadoop_data/dfs/data</value>
</property>
```

4. What is the block replication and which file is it configured in?


Block replication means no of copies of a data for fault tolerance

it is configured in "hdfs-site.xml" file

```
<property>
  <name>dfs.replication</name>
  <value>1</value>
</property>
```

Perform

1. Start HDFS and verify that it's running



2. Create a new directory /exercise1 on HDFS

```
Terminal
mohtashim@masternode: ~
mohtashim@masternode:~$ start-all.sh
This script is deprecated. Instead use start-dfs.sh and start-yarn.sh
21/06/02 21:38:40 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Starting namenodes on [localhost]
localhost: starting namenode, logging to /opt/hadoop-2.6.0/logs/hadoop-mohtashim-namenode-masternode.out
localhost: starting datanode, logging to /opt/hadoop-2.6.0/logs/hadoop-mohtashim-datanode-masternode.out
Starting secondary namenodes [0.0.0.0]
21/06/02 21:39:05 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Starting yarn daemons
Starting resourcemanager, logging to /opt/hadoop-2.6.0/logs/yarn-mohtashim-resourcemanager-masternode.out
localhost: starting nodemanager, logging to /opt/hadoop-2.6.0/logs/yarn-mohtashim-nodemanager-masternode.out
mohtashim@masternode:~$ jps
2548 DataNode
3031 NodeManager
2426 NameNode
2907 ResourceManager
2699 SecondaryNameNode
3246 Jps
mohtashim@masternode:~$ hdfs dfs -mkdir /exercise1
21/06/02 21:40:09 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
mohtashim@masternode:~$ hdfs dfs -ls /
21/06/02 21:40:24 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
drwxr-xr-x - mohtashim supergroup          0 2021-05-28 16:42 /MyDir
drwxr-xr-x - mohtashim supergroup          0 2021-06-02 21:40 /exercise1
mohtashim@masternode:~$
```

3. Upload GitHub repo sample_data/deckofcards.txt to HDFS under /exercise1 directory

```
Terminal
mohtashim@masternode: ~
mohtashim@masternode:~$ jps
2548 DataNode
3031 NodeManager
2426 NameNode
2907 ResourceManager
2699 SecondaryNameNode
3246 Jps
mohtashim@masternode:~$ hdfs dfs -mkdir /exercise1
21/06/02 21:40:09 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
mohtashim@masternode:~$ hdfs dfs -ls /
21/06/02 21:40:24 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
drwxr-xr-x - mohtashim supergroup          0 2021-05-28 16:42 /MyDir
drwxr-xr-x - mohtashim supergroup          0 2021-06-02 21:40 /exercise1
mohtashim@masternode:~$ sudo ifconfig eth0 up
[sudo] password for mohtashim:
mohtashim@masternode:~$ sudo dhclient eth0
mohtashim@masternode:~$ gedit

(gedit:3674): GLib-GObject-CRITICAL **: g_value_set_object: assertion 'G_VALUE_HOLDS_OBJECT (value)' failed
(gedit:3674): Gtk-CRITICAL **: gtk_widget_captured_event: assertion 'WIDGET_REALIZED_FOR_EVENT (widget, event)' failed
mohtashim@masternode:~$ hdfs dfs -put sample_data/deckofcards.txt /exercise1
21/06/02 21:50:50 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
mohtashim@masternode:~$ hdfs dfs -ls /exercise1
21/06/02 21:51:11 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
-rw-r--r-- 1 mohtashim supergroup          693 2021-06-02 21:50 /exercise1/deckofcards.txt
mohtashim@masternode:~$
```

4. View the content of the /exercise1 directory

```
Terminal
mohtashim@masternode: ~
mohtashim@masternode:~$ jps
2548 DataNode
3031 NodeManager
2426 NameNode
2907 ResourceManager
2699 SecondaryNameNode
3246 Jps
mohtashim@masternode:~$ hdfs dfs -mkdir /exercise1
21/06/02 21:40:09 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
mohtashim@masternode:~$ hdfs dfs -ls /
21/06/02 21:40:24 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
drwxr-xr-x - mohtashim supergroup          0 2021-05-28 16:42 /MyDir
drwxr-xr-x - mohtashim supergroup          0 2021-06-02 21:40 /exercise1
mohtashim@masternode:~$ sudo ifconfig eth0 up
[sudo] password for mohtashim:
mohtashim@masternode:~$ sudo dhclient eth0
mohtashim@masternode:~$ gedit

(gedit:3674): GLib-GObject-CRITICAL **: g_value_set_object: assertion 'G_VALUE_HOLDS_OBJECT (value)' failed
(gedit:3674): Gtk-CRITICAL **: gtk_widget_captured_event: assertion 'WIDGET_REALIZED_FOR_EVENT (widget, event)' failed
mohtashim@masternode:~$ hdfs dfs -put sample_data/deckofcards.txt /exercise1
21/06/02 21:50:50 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
mohtashim@masternode:~$ hdfs dfs -ls /exercise1
21/06/02 21:51:11 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
-rw-r--r-- 1 mohtashim supergroup          693 2021-06-02 21:50 /exercise1/deckofcards.txt
mohtashim@masternode:~$
```

5. Determine the size of the hamlet.txt file in KB that resides on HDFS (not local directory)

```
Terminal
mohtashim@masternode: ~
21/06/02 21:40:24 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
drwxr-xr-x - mohtashim supergroup 0 2021-05-28 16:42 /MyDir
drwxr-xr-x - mohtashim supergroup 0 2021-06-02 21:40 /exercise1
mohtashim@masternode:~$ sudo ifconfig eth0 up
[sudo] password for mohtashim:
mohtashim@masternode:~$ sudo dhclient eth0
mohtashim@masternode:~$ gedit
(gedit:3674): GLib-GObject-CRITICAL **: g_value_set_object: assertion 'G_VALUE_HOLDS_OBJECT (value)' failed
(gedit:3674): Gtk-CRITICAL **: gtk_widget_captured_event: assertion 'WIDGET_REALIZED_FOR_EVENT (widget, event)' failed
mohtashim@masternode:~$ hdfs dfs -put sample_data/deckofcards.txt /exercise1
21/06/02 21:50:50 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
mohtashim@masternode:~$ hdfs dfs -ls /exercise1
21/06/02 21:51:11 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
-rw-r--r-- 1 mohtashim supergroup 693 2021-06-02 21:50 /exercise1/deckofcards.txt
mohtashim@masternode:~$ hdfs dfs -ls /
21/06/02 21:53:42 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
drwxr-xr-x - mohtashim supergroup 0 2021-05-28 16:42 /MyDir
drwxr-xr-x - mohtashim supergroup 0 2021-06-02 21:50 /exercise1
mohtashim@masternode:~$ hdfs dfs -ls /exercise1
21/06/02 21:54:03 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
-rw-r--r-- 1 mohtashim supergroup 693 2021-06-02 21:50 /exercise1/deckofcards.txt
mohtashim@masternode:~$
```

6. Print the first 25 lines to the screen from deckofcards.txt on HDFS

```
Terminal
mohtashim@masternode: ~
BLACK|SPADE|10
BLACK|SPADE|J
mohtashim@masternode:~$ hdfs dfs -cat /exercise1/deckofcards.txt | head -n 25
21/06/02 22:29:09 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
BLACK|SPADE|2
BLACK|SPADE|3
BLACK|SPADE|4
BLACK|SPADE|5
BLACK|SPADE|6
BLACK|SPADE|7
BLACK|SPADE|8
BLACK|SPADE|9
BLACK|SPADE|10
BLACK|SPADE|J
BLACK|SPADE|Q
BLACK|SPADE|K
BLACK|SPADE|A
BLACK|CLUB|2
BLACK|CLUB|3
BLACK|CLUB|4
BLACK|CLUB|5
BLACK|CLUB|6
BLACK|CLUB|7
BLACK|CLUB|8
BLACK|CLUB|9
BLACK|CLUB|10
BLACK|CLUB|J
BLACK|CLUB|Q
BLACK|CLUB|K
mohtashim@masternode:~$ ^C
mohtashim@masternode:~$
```

7. Copy deckofcards.txt to deckofcardsCopy.txt

```
Terminal
mohtashim@masternode: ~
BLACK|SPADE|6
BLACK|SPADE|7
BLACK|SPADE|8
BLACK|SPADE|9
BLACK|SPADE|10
BLACK|SPADE|J
BLACK|SPADE|Q
BLACK|SPADE|K
BLACK|SPADE|A
BLACK|CLUB|2
BLACK|CLUB|3
BLACK|CLUB|4
BLACK|CLUB|5
BLACK|CLUB|6
BLACK|CLUB|7
BLACK|CLUB|8
BLACK|CLUB|9
BLACK|CLUB|10
BLACK|CLUB|J
BLACK|CLUB|Q
BLACK|CLUB|K
mohtashim@masternode:~$ ^C
mohtashim@masternode:~$ hdfs dfs -cp /exercise1/deckofcards.txt /exercise1/deckofcardsCopy.txt
21/06/02 22:30:24 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
mohtashim@masternode:~$ hdfs dfs -ls /exercise1
21/06/02 22:30:34 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r-- 1 mohtashim supergroup 693 2021-06-02 21:50 /exercise1/deckofcards.txt
-rw-r--r-- 1 mohtashim supergroup 693 2021-06-02 22:30 /exercise1/deckofcardsCopy.txt
mohtashim@masternode:~$
```

8. Copy deckofcards.txt back to local file system and name it deckofcards.copy.txt

```

Terminal
mohtashim@masternode: ~
BLACK|SPADE|J
BLACK|SPADE|Q
BLACK|SPADE|K
BLACK|SPADE|A
BLACK|CLUB|2
BLACK|CLUB|3
BLACK|CLUB|4
BLACK|CLUB|5
BLACK|CLUB|6
BLACK|CLUB|7
BLACK|CLUB|8
BLACK|CLUB|9
BLACK|CLUB|10
BLACK|CLUB|J
BLACK|CLUB|K
mohtashim@masternode:~$ ^C
mohtashim@masternode:~$ hdfs dfs -cp /exercise1/deckofcards.txt /exercise1/deckofcardsCopy.txt
21/06/02 22:30:24 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
mohtashim@masternode:~$ hdfs dfs -ls /exercise1
21/06/02 22:30:34 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r-- 1 mohtashim supergroup 693 2021-06-02 21:50 /exercise1/deckofcards.txt
-rw-r--r-- 1 mohtashim supergroup 693 2021-06-02 22:30 /exercise1/deckofcardsCopy.txt
mohtashim@masternode:~$ hdfs dfs -get /exercise1/deckofcards.txt sample_data/deckofcards.copy.txt
21/06/02 22:32:03 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
mohtashim@masternode:~$ ls sample_data
deckofcards.copy.txt deckofcards.txt
mohtashim@masternode:~$

```

9. Check the entire filesystem for inconsistencies/problems

```

Terminal
mohtashim@masternode: ~
-libjars <comma separated list of jars> specify comma separated jar files to include in the classpath.
-archives <comma separated list of archives> specify comma separated archives to be unarchived on the compute machines.
The general command line syntax is
bin/hadoop command [genericOptions] [commandOptions]
mohtashim@masternode:~$ hdfs fsck /
21/06/02 22:36:37 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Connecting to namenode via http://localhost:50070
FSCK started by mohtashim (auth:SIMPLE) from /127.0.0.1 for path / at Wed Jun 02 22:36:39 IST 2021
..Status: HEALTHY
Total size: 1415 B
Total dirs: 3
Total files: 3
Total symlinks: 0
Total blocks (validated): 3 (avg. block size 471 B)
Minimally replicated blocks: 3 (100.0 %)
Over-replicated blocks: 0 (0.0 %)
Under-replicated blocks: 0 (0.0 %)
Mis-replicated blocks: 0 (0.0 %)
Default replication factor: 1
Average block replication: 1.0
Corrupt blocks: 0
Missing replicas: 0 (0.0 %)
Number of data-nodes: 1
Number of racks: 1
FSCK ended at Wed Jun 02 22:36:39 IST 2021 in 27 milliseconds
The filesystem under path '/' is HEALTHY
mohtashim@masternode:~$

```

10. Delete deckofcards.txt from HDFS

```

Terminal
mohtashim@masternode: ~
21/06/02 22:37:42 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 0 minutes, Emptier interval = 0 minutes.
Deleted /exercise1/deckofcards.txt
mohtashim@masternode:~$ hdfs dfs -rm /exercise1/deckofcards.txt
21/06/02 22:38:00 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
rm: '/exercise1/deckofcards.txt': No such file or directory
mohtashim@masternode:~$ hdfs dfs -ls /
21/06/02 22:38:21 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
drwxr-xr-x - mohtashim supergroup 0 2021-05-28 16:42 /MyDir
drwxr-xr-x - mohtashim supergroup 0 2021-06-02 22:37 /exercise1
mohtashim@masternode:~$ hdfs dfs -ls /exercise1
21/06/02 22:38:34 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
-rw-r--r-- 1 mohtashim supergroup 693 2021-06-02 22:30 /exercise1/deckofcardsCopy.txt
mohtashim@masternode:~$ hdfs dfs -cp /exercise1/deckofcardsCopy.txt /exercise1/deckofcards.txt
21/06/02 22:39:25 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
mohtashim@masternode:~$ hdfs dfs -ls /exercise1
21/06/02 22:39:31 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r-- 1 mohtashim supergroup 693 2021-06-02 22:39 /exercise1/deckofcards.txt
-rw-r--r-- 1 mohtashim supergroup 693 2021-06-02 22:30 /exercise1/deckofcardsCopy.txt
mohtashim@masternode:~$ hdfs dfs -rm /exercise1/deckofcards.txt
21/06/02 22:39:41 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
21/06/02 22:39:42 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 0 minutes, Emptier interval = 0 minutes.
Deleted /exercise1/deckofcards.txt
mohtashim@masternode:~$

```

11. Delete the /exercise1 directory from HDFS

```
Terminal
mohtashim@masternode: ~
-r, --verbose      similar to 'rmdir a/b/c a/b a'
--help            output a diagnostic for every directory processed
--help            display this help and exit
--version          output version information and exit

Report rmdir bugs to bug-coreutils@gnu.org
GNU coreutils home page: <http://www.gnu.org/software/coreutils/>
General help using GNU software: <http://www.gnu.org/gethelp/>
For complete documentation, run: info coreutils 'rmdir invocation'
mohtashim@masternode:~$ hdfs dfs -rmdir /exercise1 | -v
-v: command not found
21/06/02 22:41:09 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
rmdir: '/exercise1': Directory is not empty
mohtashim@masternode:~$ hdfs dfs -rmr /exercise1
rmr: DEPRECATED: Please use 'rm -r' instead.
21/06/02 22:42:42 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Deleted /exercise1
21/06/02 22:42:43 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 0 minutes, Emptier interval = 0 minutes.
mohtashim@masternode:~$ hdfs dfs -rmr /
rmr: DEPRECATED: Please use 'rm -r' instead.
21/06/02 22:42:58 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
rmr: '/': Input/output error
21/06/02 22:42:59 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 0 minutes, Emptier interval = 0 minutes.
mohtashim@masternode:~$ hdfs dfs -ls /
21/06/02 22:43:14 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
drwxr-xr-x  - mohtashim supergroup          0 2021-05-28 16:42 /MyDir
mohtashim@masternode:~$
```

12. Take a second to look at other available shell options.